

Understanding FastAPI Lifespan: A Comprehensive Guide

Introduction

In FastAPI, the "lifespan" feature provides a powerful way to manage application lifecycle events - specifically what happens when your application starts up and shuts down. Think of it like the birth and death of your application, with everything that needs to happen in between.

A helpful analogy is to think of lifespan as setting up and cleaning up a workshop. Before you start working (handling requests), you need to set up your tools and materials. When you're done, you need to clean up and put everything away properly.

Why Use Lifespan?

Lifespan management is crucial for:

1. **Resource initialization:** Setting up database connections, external API clients, or in-memory caches
2. **Resource cleanup:** Properly closing connections and freeing resources when shutting down
3. **Background tasks:** Starting and stopping background processing tasks
4. **Configuration loading:** Loading environment variables or configuration files

Basic Lifespan Implementation

FastAPI implements lifespan using Python's context manager pattern. Here's the basic structure:

```
from contextlib import asynccontextmanager
from fastapi import FastAPI

@asynccontextmanager
```

```

async def lifespan(app: FastAPI):
    # Startup code: runs before the application starts accepting requests
    print("Application is starting up")

    yield # FastAPI serves requests during this yield

    # Shutdown code: runs when the application is shutting down
    print("Application is shutting down")

app = FastAPI(lifespan=lifespan)

@app.get("/")
async def root():
    return {"message": "Hello World"}

```

The code before `yield` executes during startup, and the code after `yield` executes during shutdown.

Example 1: Database Connection Management

One of the most common use cases for lifespan is managing database connections:

```

from contextlib import asynccontextmanager
from fastapi import FastAPI, Depends
from sqlalchemy import create_engine
from sqlalchemy.ext.declarative import declarative_base
from sqlalchemy.orm import sessionmaker, Session

# Database setup
SQLALCHEMY_DATABASE_URL = "sqlite:///./test.db"
engine = None
SessionLocal = None
Base = declarative_base()

class User(Base):
    __tablename__ = "users"

    # Define your model fields here

```

```

# ...

@asynccontextmanager
async def lifespan(app: FastAPI):
    # Startup: initialize database connection
    global engine, SessionLocal
    print("Creating database connection pool...")
    engine = create_engine(SQLALCHEMY_DATABASE_URL)
    SessionLocal = sessionmaker(autocommit=False, autoflush=False, bind=
engine)

    # Create tables
    Base.metadata.create_all(bind=engine)

    yield

    # Shutdown: close database connection
    print("Closing database connection pool...")
    if engine:
        engine.dispose()

app = FastAPI(lifespan=lifespan)

# Dependency to get DB session
def get_db():
    db = SessionLocal()
    try:
        yield db
    finally:
        db.close()

@app.get("/users/")
def read_users(db: Session = Depends(get_db)):
    # Use db session to query users
    return {"message": "Database connection is working"}

```

This example:

- Initializes SQLAlchemy connection at startup

- Creates database tables if they don't exist
- Properly disposes of the connection pool at shutdown
- Provides a dependency for route handlers to access the database

Example 2: External API Client and Cache Management

Managing multiple external resources:

```
from contextlib import asynccontextmanager
from fastapi import FastAPI
import httpx
import redis.asyncio as redis
import os

@asynccontextmanager
async def lifespan(app: FastAPI):
    # Startup: initialize HTTP client and Redis connection
    print("Initializing application resources...")

    # Create and store HTTP client in app state
    app.state.http_client = httpx.AsyncClient(timeout=30.0)

    # Create Redis connection for caching
    app.state.redis = await redis.from_url(
        os.getenv("REDIS_URL", "redis://localhost:6379/0")
    )

    yield

    # Shutdown: close connections gracefully
    print("Cleaning up application resources...")
    await app.state.http_client.aclose()
    await app.state.redis.close()

app = FastAPI(lifespan=lifespan)

@app.get("/external-data/")
```

```

async def get_external_data():
    # Use the HTTP client from app state
    response = await app.state.http_client.get("https://jsonplaceholder.typic
ode.com/posts/1")
    data = response.json()

    # Cache the result in Redis
    await app.state.redis.set("cached_post", str(data), ex=3600)

    return {"data": data}

@app.get("/cached-data/")
async def get_cached_data():
    # Get data from cache
    cached_data = await app.state.redis.get("cached_post")
    if cached_data:
        return {"source": "cache", "data": cached_data.decode()}
    return {"source": "none", "data": None}

```

This example:

- Initializes both an HTTP client and a Redis connection
- Stores them in the app's state for use in route handlers
- Properly closes both connections during shutdown

Example 3: Background Task Management

Managing background processes:

```

from contextlib import asynccontextmanager
from fastapi import FastAPI
import asyncio
import time
from datetime import datetime

@asynccontextmanager
async def lifespan(app: FastAPI):
    # Create a shared variable to control the background task
    app.state.should_exit = False

```

```

# Define the background task
async def background_task():
    while not app.state.should_exit:
        print(f"Background task running at {datetime.now()}")
        app.state.last_run = datetime.now()
        await asyncio.sleep(10) # Run every 10 seconds

# Start the background task
task = asyncio.create_task(background_task())
app.state.background_task = task

print("Application started, background task is running")
yield

# Signal the background task to exit
print("Shutting down background task...")
app.state.should_exit = True

# Wait for the task to complete (with timeout)
try:
    await asyncio.wait_for(app.state.background_task, timeout=5.0)
except asyncio.TimeoutError:
    print("Background task did not exit in time, cancelling...")

print("Application shutdown complete")

app = FastAPI(lifespan=lifespan)

@app.get("/background-status")
async def get_status():
    if hasattr(app.state, "last_run"):
        last_run = app.state.last_run
        return {"status": "running", "last_run": last_run}
    return {"status": "not started"}

```

This example:

- Creates and manages a background task that runs periodically

- Stores the task and its state in the app's state object
- Properly handles cleanup during shutdown with timeout protection

Best Practices for Using Lifespan

1. **Keep startup code efficient:** Slow startup will delay application readiness
2. **Handle exceptions gracefully:** Use try/except in both startup and shutdown code
3. **Set reasonable timeouts:** Don't let shutdown tasks hang indefinitely
4. **Use app.state for shared resources:** Store connections and shared state in app.state
5. **Log critical events:** Log startup/shutdown progress for debugging

Lifespan vs. Event Handlers (Deprecated Method)

In earlier versions of FastAPI, lifecycle events were managed with `@app.on_event()` decorators:

```
# Old way (still works but deprecated)
@app.on_event("startup")
async def startup_event():
    # Startup code
    pass

@app.on_event("shutdown")
async def shutdown_event():
    # Shutdown code
    pass
```

The newer lifespan approach is preferred because:

- It guarantees that shutdown code runs even if startup fails
- It follows Python's context manager pattern, which is more idiomatic
- It's more maintainable since related startup/shutdown code is in the same function

Conclusion

FastAPI's lifespan feature provides a robust way to manage application lifecycle events. By properly initializing and cleaning up resources, you can build more reliable and efficient web services that handle both expected and unexpected termination gracefully.

Remember to think of lifespan management as preparing your workshop before you start working, and cleaning up properly when you're done. This ensures your application runs efficiently and doesn't leave any resources hanging when it shuts down.